

Capstone Project

Radek Chramosil

WorldQuant University

E-mail: radek@keemail.me

This document is the Capstone Project Report. It details a simple database structure to handle the second market order book data, shows a simple strategy using the data set and it discusses many topics in the high frequency market.

Keywords— High frequency trading, Market microstructure, Algorithmic trading, Market order book modelling

Introduction

The asset is Bitcoin quoted in USD. The access to the full exchange is free and level 2 data is available for every tick. The data was collected on second basis for one hour. The collection was done with 1 second frequency as opposed to per trade basis. The sample data set is therefore some 3440 full market book states; some 100 quotes for bid prices and some 100 quotes for ask prices for each state.

This documents further focuses on a general electronic trading chain, benefits and problems with high frequency trading, direct market access (DMA) with its risks. Further, I propose a simple mathematical supply and demand model for market order book and use this model for a simple trading strategy and evaluate the strategy.

High Frequency Trading Discussion

This sections highlights some issues in the so called capital market micro-structure. I focus on technical aspects mostly. It is often difficult to identify or read about the specifics, say, of an API as used by an official exchange. So, I based my work on Bitcoin and Binance exchange as an example. The discussions is focused on US stock market as this is the most commonly covered topic in the literature. This is continuous action order-driven market. I provide the time units definitions for convenience.

Second to lower time intervals	
1 second	1 000 milliseconds
1 second	1 000 000 microseconds
1 second	1 000 000 000 nanoseconds

Electronic Trading Chain

There are many aspects that can be discussed with regards to trading chain links between brokers, exchanges, investors (retail, institutional or algorithmic trading engines). However, Durbin (5) mentions that for a high frequency firm access the exchanges for two reasons only. One is to submit their trade flow (trade execution) and the other is the receipt of market data. This can be generalised to any market participant. Hence, I provide the Figure 1 that shows a high level overview.

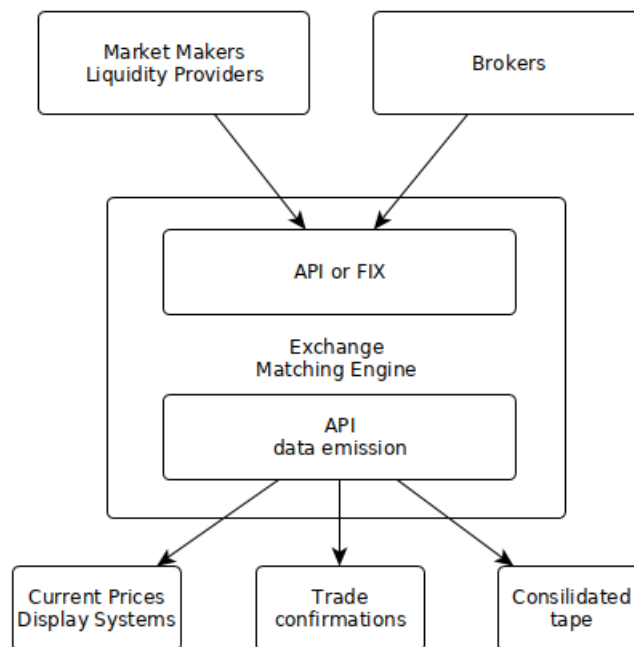


Figure 1: Market Structure Overview

Matching Engine The core about any trade chain the matching mechanism. Johnson (10) provides high level structure for a matching engine. I highlight it here to give a notion for the coding structure of a matching engine. Each order (that is price and quantity for bid or ask) is processed as it arrives to the matching engine. The below steps are taken each time there is a new order.

- The order instruction is added to internal order book
- Trade matching rules are then used to see if any matches are now possible
- The order book is updated to reflect any changes

- Execution notifications are sent to any result matches

These steps essentially are the core of any order driven market. Figure 2 shows how LSE displays their market order book to their clients (15).

SETS screen

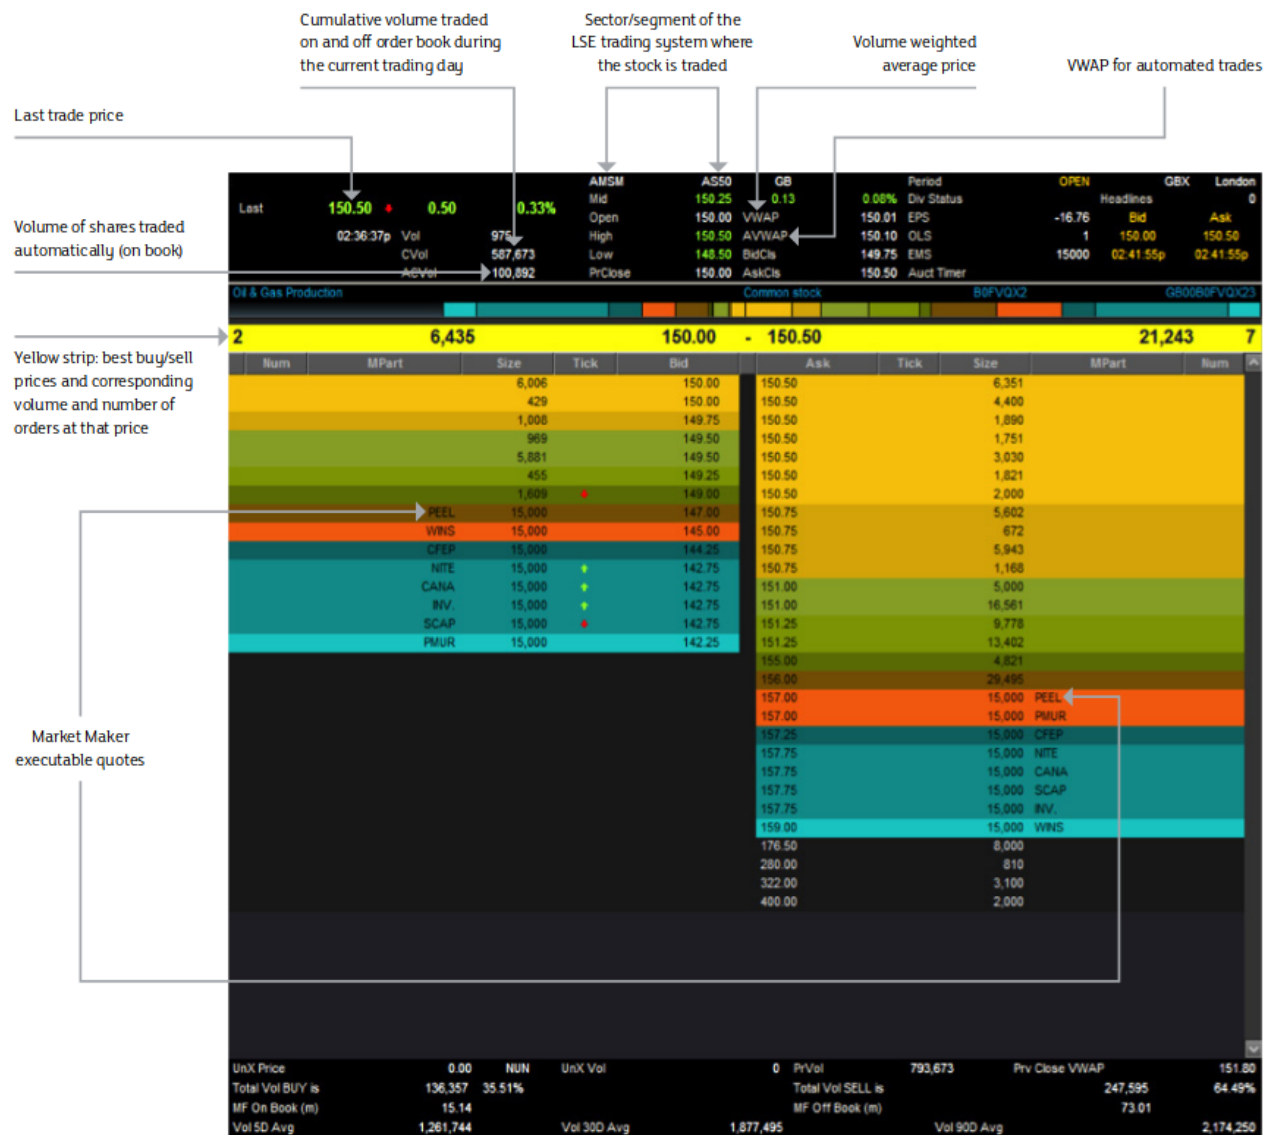


Figure 2: LSE Market Order Book screen

In an attempt to assess how this matching algorithm might work, I used Binance exchange as an example. This exchange sees the order book at a new state each time there is an order and uses a data

field 'UpdatedID' to mark this state. The exchange members are then able to request the order-book state stamped by this ID. Binance uses only REST API which we can query programmatically, say using Python's Requests package. On the contrary, on the NYSE's website we can read about different services this exchange offers to their clients (16). We see that they have an event-based depth of book feed available, which is not something Binance offers as yet. Programmatically, event driven API would be a better option for us, however we can expect that NYSE charges their fee for this service.

To understand how to classify the different levels of service an exchange might provide, it is useful to consider the latency differences. O'Hara (18) provides a good summary of services for Tokyo exchange. This exchange offers Arrownet, the standard service, which has the latency of several milliseconds. Their priority service allows users to place devices at data center entry points for their network, lowering the latency to 260 microseconds. Finally, the co-location at their primary side offers 15.7 nanoseconds service. We can learn from this that the simplest and most straight forwards way to classify the services the exchange provide is based around the time latency. Unfortunately, the exchanges are not transparent around the latencies. NYSE does not talk about latencies at all on their website (16). Binance simply leaves this issue to the market participants as any one can access their matching engine over the Internet.

Further, the important aspect of co-location is that once co-located the market participants have the cable of the same length, see the interview with Manoj Narang (21). Once, we pay for the co-location service, we should have the same delay as anybody else who pays for the service.

Financial Information Exchange protocol (FIX protocol, see Johnson (10) for details, it is just a version of XML) and exchange application programming interface (API) provide a way to submit orders. The broker orders as well as market makers therefore must reach the matching engine the same way even though that the difference is that some people press buttons on their screen and some people monitor the market programmatically and their programme will submit the order. FIX should be slower as it needs to be interpreted as opposed to API but it is used as it is common to the whole market. Johnson (10) also mentions exchanges have separate servers for processing orders and emitting data. We would have to investigate what a specific exchange does before we start trading.

After a trade is submitted to the matching engine, the matching engine must make its state available to the market, specifically market maker, broker, data vendors or authorities. This includes trade confirmation but also open orders as well as the matched trades. In the US an exchange may submit the confirmation firstly to Consolidated Tape Authority (CTA) - see Durbin (5). CTA then in turn should have all market data and some market participants use it as a data source (Durbin (5) mentions hedge funds). While Europe is missing the consolidated type all together as suggested by O'Hara (18). Perhaps, the rest of the world faces this issue too.

Brokers The traditional way for people to access the market is over the phone via a broker; refer to Harris (6) to see what specific applications or systems are involved. I just mention that in the principle for the order-driven market the order must end up in the same order book. However, brokers actually route the market order to the venue that pays them the most, see Harris (6) or O'Hara (18). O'Hara mentions that there are some 15 different stock exchanges and 50 alternative venues.

In order for the broker to give their clients competitive edge over poor algorithmic traders, brokers now offer algorithmic execution for trades as well. As an example, I look into Investment Technology Group's (ITG) execution strategies (8). Their clients can benefit from a large variety of different execution algorithms or the ability to build their own. ITG also offers DMA in some forms. So, we see that

the playing field is levelling up even for any market player.

Algorithmic Trading Engines The structure for an algorithmic trading engine is reasonably well described by Durbin (5). He provides summarises his ideas in the below flow chart in Figure 3.




Figure 3: High Frequency Trader - Operational Structure

He defines the terms in the following. *Thinker* takes direction from humans. This would be the graphical interface for the application. Then again we have *exchange*. This is the matching engine as discussed above. *Listener* takes in the market data and make them available to other components. Durbin states that this should sit as close as possible the market data source (exchange or CTA) and take in every tick. Listener also does checks for errors or arbitrage if multiple sources disagree. Durbin then states that *Pricer* component calculates the price of assets in the real-time, however we know that we could use any strategy or use any mathematical model here, not just pricing. This component or application receives the inventory information from *Listener*. *Trader* submits/cancels the orders in the exchange. Finally, *Manager* is a program that controls all the other components.

We can see how a general algorithmic trading operation may run from Durbin's discuss, however we know that the reality of algorithmic trading is more complicated than that. We also see that this structure fits well into our discussion well, see Figure 1. Further, Johnson (10) provides code C++ snippets that can be established in individual components.

DMA Services Johnson (10) defines DMA as a computerised system responsible for executing the orders to buy or sell a given asset, rather than being worked manually by a trader. It should be noted that

the broker can give their client this access rather than the exchange itself, however many exchanges give DMA to their market makers.

Benefits and Pitfalls of DMA technologies

DMA Types Johnson (10) provides a good overview of DMA services available. I just provide a brief definitions and suggest a way how we should study the pros and cons as market participants.

Direct Market Access In section, we define *Direct Market Access* (DMA) as the situation when broker's client makes use of the broker's infrastructure to send the orders to the exchange, like broker's own orders. *Sponsored Access* term is used when the broker's client is given broker's access credentials to execute the trades directly at the exchange, but not having to use the broker's infrastructure. This should be faster, however there still is some connection overhead for the broker to monitor the trading to prevent any excessive risk taking by the client. *Crossing* refers to the situation when brokers offers the access to the dark pool to their clients. This, Johnson claims, to be used when a client is looking to carry out a large trade and is trying to hide it from a many market participants that are able to respond to the orders visible in exchange order book. *Direct Liquidity Access* is then the term when the broker offers the access to dark pools and to the exchanges at the same time. In addition, some brokers started to use the term *Direct Strategy Access*. This means broker provides their client with execution algorithms.

Thus, we see that the way to compare different ways to access the markets can be classified by

- latency,
- venue availability,
- execution algorithm (that is how the trade is routed and how a single large trade is distributed to smaller trades),
- price.

We would need to research what specific deals are available to us from different brokers and decided on the best deals based on the above criteria.

One final thought about DMA comparison is that in case of crypto-currency market, there are no brokers at all. One additional form of DMA is not have a broker.

As for the benefits, the obvious case for them that automated markets are more efficient so cheaper, that is the trading is cheaper, which was the original cause for automation, see Patterson (19). This authors talks about the fact that NASDAQ market maker kept their spreads artificially high in the past. The solution was to automate the trading. Also, one more obvious benefits is that direct market access provides greater transparency with regards to the price data as well as market order book. The exchange members can request all the data their desire, however in case of standard capital markets this subscription is paid. DMA technologies may allows us to bypass brokers, who often have a lot conflicts of interest with respect to their clients (broker can sell their client order flow or they execute market orders on behalf of their clients when they deem that the position is too risky and collect the trade fee). I test DMA and a trading strategies under live access to a crypto-currency exchange.

On the other hand, the market automation created problems at the same as it had provided us with benefits. Exchanges now sell access to a public information, they sell access to the market order book

with different latencies. As a results, they turn public information into private information for time intervals measured in microseconds or nanoseconds (18). In fact Brad Katsuyama claimed in his interview with BBC that the exchanges make more money by selling high speed data and technology than they make by matching trades . It is clear that some market participants benefit from this latency advantage at expense of the market participants that did not pay for the low latency access, see Patterson (19). Another problem is that the matching engines are not publicly available in standard markets. What this means is that the exchange can create special or hidden order types, which are again given just to some specific market participants, see Bodek (2). This author claims that his algorithmic trading operation went in the red as he did not have access to the hidden marker order types while his competitors had. He talks about hidden Alphas in different order types. So, we then have the problem that the matching engine is not public information just the API is.

Therefore, we can see that algorithmic trading creates transparency as more data is available nowadays while at the same time it does exactly the opposite in some other sense.

Risks in DMA

DMA does have obvious risks. Knight Capital lost \$440 million in trading loss due to a bug in their systems (17). The company suffered power disruptions and it was not able execute orders on some other day. AXA Rosenberg lost \$217 million in client money due to an error in their model for managing risk (4). So, we see that one of the big problems and risks with DMA is that when people get it wrong, they are likely to lose all their capital within a single day or worse.

Janecek (9) talked about the issues he encountered with his firm. Basically, we should consider a number of risks with regards to risk with DMA. We cannot have code bugs left in the system. This can be avoided by techniques such as test driven development. More over the computer network may fail due to computational overload in highly volatility markets. This may impact the whole system that we use. Again, we would avoid this by capacity testing for a the network. Moreover, we could extend this idea of an error caused by the market conditions to the Flash Crash. Only Renaissance Technologies had their systems ready to face Flash Crash (19), while Peterffy's company and Citadel were at lost (24).

Further risks are that we can have a mathematical modelling error or a parameter error. We can only test for these statistically or via optimisation using back testing with regards to some cost and benefit functions (such as a return function).

Moreover, the all technical risks that linked to running a data centres such as power shortages, car or track accidents causing evacuation of the area, flights patterns, flat plains, nuclear plants, rail lines, crime level and even light night strikes. All of these events could cause an outage in a data facility (23). DMA is directly dependent on all of these as it is carried out in the data centres physically next to where the matching engine server is located. In case of computer networks, we need to add risk that the connections may fail for the same reasons and many other as in case of a data centre.

As previously mentioned, we do not have a full access to the matching engine, so we run a risk that the exchange gives different order types to different market players.

There are also topics such as slippage, front-running, quote-staffing (see, Durbin (5)), however there are more general to the market structure not to the DMA itself.

Latency Types

The basic ideas behind the latency issue are well summarised by Kirilenko (11). He points that a standard way to measure latency is by determining the time it takes a given message to travel from source to destination and back - *round-trip time* (RTT). he talks about

- *communication latency* - can be reduced by co-location
- *market feed latency* - can be reduced by buying proprietary data feed
- *trading platform latency* - cannot be reduced by a trader, taken as given

Kirilenko and Lamacie (12) also points that the latency is a random variable as in Figure 4. We can guess what the RTT distribution properties are from these plots.




Figure 4: RTT - All Types: Submissions, Modify, Cancel - from (12)

Further, we can break down the latencies in a different way. Kumar (13) provides a useful overview of how trading system latencies can be defined and measured. He summaries the problem at hand as:

$$\begin{aligned} \text{Trading System Latency} = & (\text{All the processing node latencies}) \\ & + (\text{Operating system network stack latencies}) \\ & + (\text{The entire network device latencies}) \\ & + (\text{The entire wire time latencies}) \end{aligned}$$

Kumar (13) then defines:

- *Processing node latency* is the time spent within the application business rules and logic.
- *Operating system network stack latency* is the time spent within operating system calls and hardware network interfaces.

- *Network device latencies* is the time spent in networking equipment (load balancers, firewalls and routers).
- *Wire time* is the time it takes for a message to be transmitted on a physical medium. Specifically, a foot of a cable approximately equates to a nanosecond (1). It is 50 to 55 micro seconds to data to travel one way on 10 km of fiber optic (13).

This discussion shows that to reduce the latency within a computer system that we need to consider every step our computer does along the way between our computer and the matching engine. That we would have to use the most efficient algorithms and data structures to handle the processing node latency. The industry uses real-time operating systems (5). Alternative way to use the just a chip set without the operating system. The work documented by Algo-Logic may serve as a good example (14).

It is worth noting that a lot of low latency issues are linked to the US equity markets and the equity instruments that have multiple venues which then leads to the latency arbitrage, however one way to solve this is to just use instruments that have only one trading venue. We may stay from the race altogether. As a result, HFT firm may build its business around mathematical modelling, see how the company called RSJ does the business (22).

0.1 Typical Order Messaging System

I was not able to find any specific information with regards to type order messaging process - Order Request, Data Conversion, Order Enrichment, Order Transmission and Exchange Acknowledgement for any standard exchange. The Investor Exchange provides REST API (7) for the data but not for orders. We check how to communicate to Binance exchange using a python package (20), however that is also REST API.

Low latency architecture We are left with general suggestions about how to reduce latency. C++ is seen as the fastest language that gives control over the computer at very low level (memory management). Generally, it is good to keep the data in memory as well so many algorithmic trading operations use KDB, which is a in-memory relational time series database. The coding style should be modular and functional as opposed to object orientated as inheritance slows the program down. In addition, jobs like error handling comes with an overhead. So, less steps in the code, the better, enough though we might be short on standard coding practices. As for operating system, Durbin (5) mentioned Linux (CentOS is mentioned by Lockwood (14)). However, we still should remove a lot of functionality, such as a graphical interface, for the server where would trade from. The point is that we wish to run our application instantaneously and at the constant time. We wish to remove as much randomness from our latencies as possible from both the communication network and our system. So, we can control the system as much as possible.

Obvious options to consider are parallel processing within a single server. This is both CPU multi-threading (Durbin (5) mentions Threading Building Blocks C++ runtime library by Intel) but also GPU processing as used in TensorFlow, however in our case we would not worry about neural networks but focused on, say, pricing 1000 options at the same time. The distributed computing is not used not used as there is an overhead to distributing the computation across multiple computers.

If we wish to go more experimental, we could start looking into new operating system or coding languages. Three-year-old Redox OS is a MIT licensed micro kernel with some 20 000 lines of code

fully written in coding language called Rust. It is reasonable to expect this OS to be faster than Linux as Linux has millions of lines of code. Redox OS has modular structure for drivers, so we would be able to amend or change drivers as we like. Linux does not have this modularity. In addition, Rust is a new c-based language, it borrows from C++ and improves on the weakness of C++, for example it has native support for multi-threading. On the contrary, both Rust and Redox OS are not tested or industry proven as C++ and Linux is, however they are worth looking into.

To sum up, the high frequency space is ever changing and all we can do is to keep as best as we can or have the time to do so.

Strategy with High Frequency Data

This section details my work with DMA and high frequency data.

Database Structure

I used the below simple database structure. This was carried out to store whole market order book for a currency pair USD quoted in Bitcoin. The data was requested using the Binance's REST API. The schema for the database is shown in Figure 5.

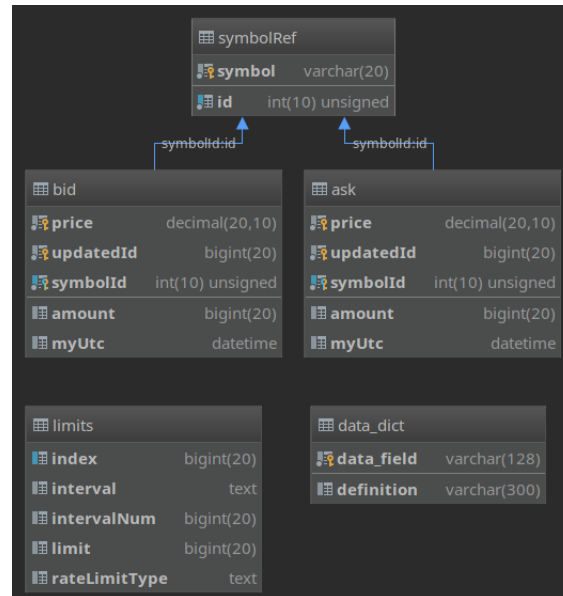


Figure 5: Schema

Additionally, I built functionality that can create tables on fly using Pandas and SQL Alchemy, however the default behaviours are not very good as it fails to create foreign keys and so on.

Tables

symbolRef This is just a reference table for symbols. The symbols are strings, however it is more efficient to use integers for foreign keys and it is better to store integers than strings.

bid and ask These two tables are identical from a data types point of view however they just store the bid and ask prices as well as my UTC time stamp and the server *'updatedId'* flag. This is a large integer that gets changed each time there is a trade.

limits . Basically, the exchange allows users to request or order limits in time intervals. We can make 1200 data requests in a minute or pass in 10 orders in a second and/or 100 000 orders in a day. If we are creating a real trading application we would have to consider these closely.

data_dict This is a manually created table and holds descriptions for the data fields in the database.

Bid and Ask

Bid and ask prices are split into two tables. The tables are identical. The tables from Figure 6 and Figure 7 carry the data for the symbol and specific Exchange ID as given by *'updateId'* column. My computer time stamps are in the *'myUtc'* column representing the time on my computer.

price *	amount *	updatedId *	myUtc	symbol *
0.00020497	1806	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020503	147	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020504	6	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020511	9	48208950	2018-11-19 21:03:04	TUSDBTC
0.0002052	4453	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020538	4453	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020546	18528	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020556	4453	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020558	1194	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020576	34	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020577	15	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020586	5	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020587	5	48208950	2018-11-19 21:03:04	TUSDBTC

Figure 6: Ask Price

price *	amount *	updatedId *	myUtc	symbol *
0.00020471	504	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020468	183	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020466	281728	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020465	10	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020463	49	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020461	20	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020456	20	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020455	10	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020454	5	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020453	10	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020452	10	48208950	2018-11-19 21:03:04	TUSDBTC
0.0002045	6	48208950	2018-11-19 21:03:04	TUSDBTC
0.00020449	30	48208950	2018-11-19 21:03:04	TUSDBTC

Figure 7: Bid Price

Exchange API

It is important to handle the communication to the Exchange well. The class '*DataEndPoint*' hold all methods that can collect the data from the exchange. I leave it to the reader to read over the comments for these methods to see what can be easily pulled out from the exchange. The import aspect of communicating with an exchange is not to breach the request or order limits, see Durbin (5).

Data Analysis - First Look

I collected an hour of trading data on per second basis. Figure 8 shows quoted amounts with bid and ask prices within one minute of trading. We see that there is a lot of volatility over a minute interval.




Figure 8: All bid and ask prices and quantities in one minute

The problem in the Figure 8 is that we do not see much details. So, I plotted the same time interval but with close look at what was happening around the spread in Figure 9. Again we see the volatility in the price within one minute.

I also consider the maximum bid and minimum ask prices and plot them over the one hour interval. This is plotted in Figure 10. This shows that we do have trends and turning points happening within a trading hour.

Finally, we check the spread for the two currencies as in Figure 11. we note the currency spread seems stationary, which is a common phenomena in financial or currency markets.




Figure 9: Detailed Price and quantity moments around the spread in one minute




Figure 10: Max bid and min ask over one hour




Figure 11: Max bid and min ask spread over one hour

Model

I tried to create my own supply and demand model as implied in the market order book. That is the model would suggest to us a trade when there is a large deviation from market equilibrium. I thought about market middle price, P_{mid} , as

$$P_{mid} = \frac{Bid_{max} + Ask_{min}}{2}$$

where Bid_{max} is the maximum bid price for a given tick and Ask_{min} is the minimum ask price for the same tick.

The idea of my supply and demand estimator should give oscillating estimates around this middle price based on the current supply and demand. I attach a probability weights to each supply and demand amounts (see Figure 8 and 9 to get some feel for what the cumulative amounts look like). I assumed that the probability of arrivals of new orders follows log-normal probability where the orders arriving close to the maximum bid and minimum ask prices are most likely and as we go away from this bid and ask spread, the arrival of new orders is less likely so they carry smaller and smaller probability weights.

Bury (3) provides a fitting procedure for log-normal distribution. We used maximum likelihood estimator for mean, μ , and an unbiased estimator for standard deviation, σ , such that

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n \ln(x_i), \quad \hat{\sigma} = \sqrt{\frac{1}{n-1} \sum_{i=1}^n [\ln(x_i) - \hat{\mu}]^2}$$

where x_i is the bid or ask price for the given state of the order book. I assumed that the bid prices and ask prices are independent and they both come from a log-normal distribution that is estimated using the current bid and ask prices separately.

We then use the relationship between log-normal and normal distribution to find the probabilities for each prices quoted in the order book. We can use the below to find the probability, P_{x_i} , of a given price level to occur.

$$P_{x_i} = F(x; \mu, \sigma) = \Phi \left(\frac{\ln(x_i) - \mu}{\sigma} \right)$$

with Φ being the standard normal cumulative density function, see Bury (3) for details. We know that the log-normal density's shape is suitable for the ask prices, however I simply reverse the order of the probability mass programmatically for the bid prices to achieve the notion that the prices close to bid are more likely to occur. I then find the expected bid and ask amounts, a , as

$$\mathbb{E}[a] = \sum_{i=1}^n P_{x_i} a_i$$

We then have two single numbers that are representing supply and demand. The expected supply and demand numbers are probability weighted towards the current bid and asks. Further, I define total supply, S_T , and total demand, D_T as the sum of the amounts for bid and ask.

As suggested previously, I tried to define an oscillator around the middle price. I used the price adjustment, PA , as

$$PA_{ask} = \frac{\frac{S_T + P_T}{2} - (\mathbb{E}[a]_{Demand} - \mathbb{E}[a]_{Supply})}{S_T + P_T} \times \left(\frac{Ask_{min}}{Bid_{max}} - 1 \right)$$

and

$$PA_{bid} = \frac{\frac{S_T + P_T}{2} - (\mathbb{E}[a]_{Supply} - \mathbb{E}[a]_{Demand})}{S_T + P_T} \times \left(\frac{Ask_{min}}{Bid_{max}} - 1 \right)$$

When $\mathbb{E}[a]_{Demand}$ is greater than $\mathbb{E}[a]_{Supply}$ I compute

$$\text{adjusted price} = P_{mid} / (1 + PA_{bid})$$

and otherwise,

$$\text{adjusted price} = P_{mid} * (1 + PA_{ask}).$$

The results of the computation are presented in Figure 12. We see that in each case my adjusted price is above or below the actual quoted bid or ask. We note that we do not have an oscillator that would change tick by tick, which is what I was aiming for. We further note that our oscillator is never within the bid and ask spread. Instead, the oscillator stays in below bid or above ask prices for minutes at the time.

As a result, it rather indicates when there is demand stronger than supply and vice versa. My methods needs further work, however I decided to use the above calculation to design a simple trend following strategy as given in the next section.




Figure 12: Supply and demand indicator

Strategy

As shown in the Figure 12 shows that we could use my model to simple follow the strong demand as an up-trend and strong supply as a down-trend. Figure 13 shows the period with frequent switches between the strong supply and strong demand. Perhaps, it indicates a turning point in the market. Then, we do not wish to follow any trend. Figure 14 and 15 then show how we can combine the regime switches to a strategy. I used one last minute (60 observations) as an indication as to whether we should take a long a short position.

Figure 16 then show the final position. We take a long position, when the market is in the up-trend or wise-versa. We take no position when supply are demand are frequently switching . The logic for that is

- Go long when demand is winning over 55 out of last 60 observations
- Go short when supply is winning over 57 out last 60 observations
- We take no position otherwise.

Figures 14 and 15 show the mirrored regime switches. We see that they are identical so we can use just one of them. The cumulative returns for this strategy are plotted in Figure 17. The Python code for this strategy is shown below.

Python loop for smoothing the regimes

```
df['position'] = 0
for n in range(len(df)):
    if n > 60:
        view = df.loc[n-60: n, 'long_signal'].sum()
        if view > 55:
            df.loc[n, 'position'] = 1
        elif view < 3:
            df.loc[n, 'position'] = -1
        else:
            df.loc[n, 'position'] = 0
# close the trade at the end of the time period
df.loc[len(df)-1, 'position'] = 0
```




Figure 13: Possible turning point




Figure 14: Short position indicator




Figure 15: Long position indicator




Figure 16: Position - Strategy




Figure 17: Position - Valuation

Strategy Summary

In Sample

The in-sample Key Performance Indicators (KPI) for simulated trading within the hour of trading are presented below. The KPIs are calculated on the middle price so it is just an approximation to real trading. No bid-ask spread or trade fees were considered.

KPIs	
annualised_std	0.1124887732
avg_loss_return	-0.0022179186
avg_trade_return	0.0004751018
avg_win_return	0.0031681222
cagr	1.0285530229
lake_ratio	0.0025649288
max_consecutive_losers	3.0000000000
max_consecutive_winners	2.0000000000
max_drawdown	0.0071235493
mean_annualised_return	0.0836999763
number_of_trades	8.0000000000
sharp_ratio	0.7440740435
win_loss_ratio	-1.4284213327
win_rate	0.5000000000

Figure 18: KPIs - Sample trading

We see that there is a potential in this trading trading strategy based on the sample.

Out of Sample

The maximum bid and minimum ask prices for out of sample data is presented in Figure 19. The data was collected for some two or three hours of trading in seconds. We note that there has been some rapid swings in the price. The strategy was not tested or prepared for this so we will see how well it can perform.

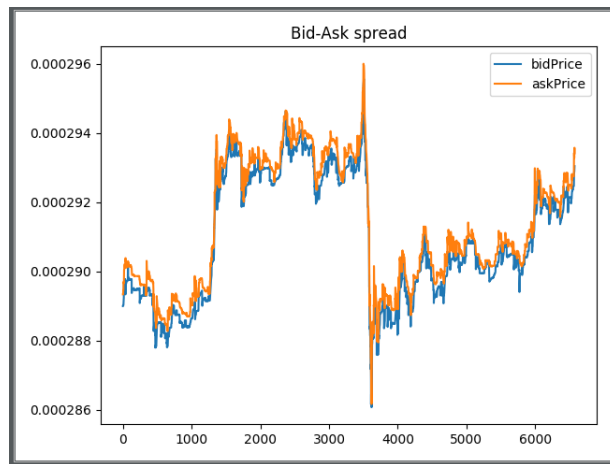


Figure 19: Out of Sample - Max Bid and Min Ask prices

Figure 20 shows the overview of the previously constructed indicator. It does not appear to be far off. Basically, the green line is when we should be long and the red line is when we should be short.

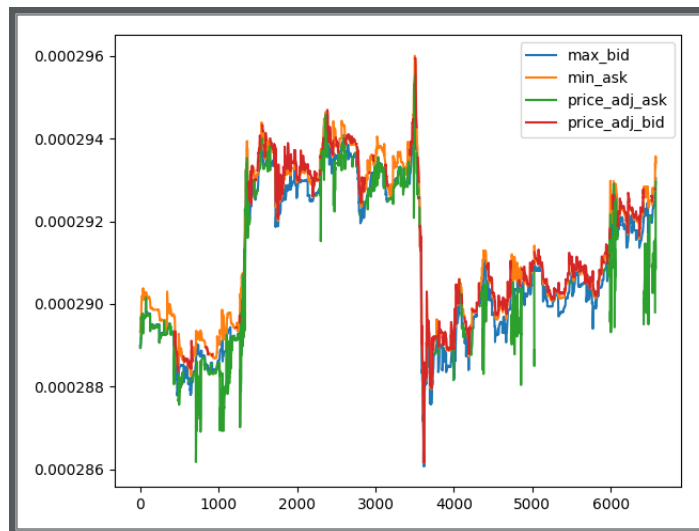


Figure 20: Out of Sample - Strategy Indicator

Figure 21 indicates the positions the trading strategy has taken out of sample.




Figure 21: Out of Sample - Trades

Figure 22 marks the out-of-sample strategy valuation as cumulative returns.




Figure 22: Out of Sample - Valuation

We see that we have made losses on the short positions while the prices was trading upwards by the end of out-of-sample testing. Further, KPIs are presented in Figure 23. We see that the win rate is quit good, however the average losses are greater than our average wins. Overall, we lost the money in the out-of-sample trading, however this approach does have the potential. We need to further investigate how to improve our approach. We should review the model used and make sure that it is statistically valid and not based on assumptions. We should review the smoothing algorithm for actually taking a view on the market. We see that we had a large jumps in the prices in the out-of-sample trading. These are not accounted for in our strategy very well.

Also, we should note that the usual KPIs do not really work very well with high frequency data. We would need to design a different set of KPIs for this purpose.

	KPIs
annualised_std	0.1039423643
avg_loss_return	-0.0021935670
avg_trade_return	0.0001943808
avg_win_return	0.0019853417
lake_ratio	0.0126090992
max_consecutive_losers	5.0000000000
max_consecutive_winners	4.0000000000
max_drawdown	0.0299282997
mean_annualised_return	-0.0591354481
number_of_trades	56.0000000000
sharp_ratio	0.5689253705
win_loss_ratio	0.9050745674
win_rate	0.5714285714

Figure 23: Out of Sample - KPIs

References and Notes

1. BBC *Slowing down the predatory high speed traders*. BBC, UK, 2017
2. Bodek, Haim *The Problem of HFT*. CreateSpace Independent Publishing Platform, US, 2013
3. Bury, Karl *Statistical Distributions in Engineering*. Cambridge University Press, UK, 1999
4. CBS News *AXA Rosenberg Paying \$242M To Settle Case On Code*. CBS News, US, 2011
5. Durbin, Michael *All about high-frequency trading*. McGraw-Hill, New York, NY, 2010
6. Harris, Larry *Trading and Exchanges: Market Microstructure for Practitioners*. Oxford University Press, Oxford, 2003
7. IEX <https://iextrading.com/developer/docs/>, last visited December 15, 2018
8. Investment Technology Group *Global Execution Strategies*. Dublin, Ireland, 2013
9. Janeček, Karel *The Importance of Failures - Talk*. CreativeMornings HQ, Prague, 2014
10. Johnson, Barry *Algorithmic trading & DMA: An introduction to direct access trading strategies*. Myeloma Press, London, UK, 2010
11. Kirilenko, Andrei *HFT Latency Regulation*. London, UK, 2016
12. Kirilenko, Andrei; Lamacie, Guilherme *Latency and Asset Prices*. London, UK, 2015
13. Kumar, Ganesh *High Frequency Trading systems and Latency measurements*. London, UK
14. Lockwood, John W., Madhu Monga *Implementing Ultra-Low-Latency Datacenter Services with Programmable Logic*. IEEE Computer Society, US, 2016
15. London Stock Exchange *A guide to London Stock Exchange trading services for equity securities*. London, UK, 2015
16. NYSE <https://www.nyse.com/market-data/real-time>, last visited December 15, 2018
17. NY Times *Knight Capital Suffers Power Failure*. NY Times, US, 2012
18. O'Hara, Maureen *High frequency market microstructure*. Journal of Financial Economics, NY, 2014
19. Patterson, Scott *Dark Pools: The Rise of the Machine Traders and the Rigging of the U.S. Stock Market*. Crown Business, London, UK, 2013
20. Python-Binance <https://python-binance.readthedocs.io/en/latest/index.html>, last visited December 15, 2018
21. Reuters *HFT Debate With Haim Bodek and Manoj Narang*. Reuters, US, 2014

22. RSJ <https://www.rsj.com/en/securities/about-us>, last visited December 15, 2018
23. VPRO *Flash Crash 2010 - VPRO documentary*. VPRO, US, 2012
24. VPRO *The Wall Street Code - VPRO documentary - 2013*. VPRO, US, 2013